CrocoLakeTools: A Python package to convert ocean observations to the parquet format

Enrico Milanese ©¹, **David Nicholson** ©¹, **Gaël Forget**², **and Susan Wijffels**¹

1 Woods Hole Oceanographic Institution, United States of America **2** Massachusetts Institute of Technology, United States of America

Summary

Investigations of the ocean state are possible thanks to the ever growing number of measurements performed with multiple instruments by different research missions. The vast and variegated efforts have brought the community to define data storage conventions (e.g. CF-netCDF) and to assemble collections of datasets (e.g. the World Ocean Database). Yet, accessing these datasets often requires the usage of multiple tools, is inefficient over the cloud, and presents an overall high entrance barrier in terms of knowledge and time required to effectively access these resources. CrocoLakeTools is a Python package that addresses those shortcomings by providing workflows to convert several datasets from their original format to a uniform parquet dataset with a shared schema.

Statement of need

`CrocoLakeTools` is a Python package to build workflows that convert ocean observations from different formats (e.g. netCDF, CSV) to parquet. CrocoLakeTools takes advantage of Python's well-established and growing ecosystem of open tools: it uses `dask`'s parallel computing capabilities to convert multiple files at once and to handle larger-than-memory data. `dask` is already well-integrated with `xarray` and `pandas`, two widely used Python libraries for the treatment of array and tabular data, respectively, and with `pyarrow`, the API to the Apache Arrow library which is used to generate the parquet dataset.

Parquet is a data storage format for big tidy data which presents several advantages: it is language agnostic (it can be accessed with Python, Matlab, Julia and web developement technologies); it offers faster reading performances than other tabular formats such as CSV; it is optimized for cloud systems storage and operations; it is widespread in the data science community, leading to a multitude of freely accessible tools and educational material.

`CrocoLakeTools` was developed with the goal of building and serving CrocoLake, a regularly refreshed database of oceanographic observations that are pre-filtered to contain only quality-controlled measurements. `CrocoLakeTools` was designed to be used by researchers, engineers and data scientists in oceanograhy and to be accessed by the wider oceanographic community.

State of the field

The need for homogenized ocean data product has prompted several efforts.

The most comprehensive project to date is likely the World Ocean Database (WOD), which contains more than 40 variables gathered from several sources (buoys, gliders, floats, bathythermographs, etc.). The earliest record dates back to 1772, the data is quality-controlled, and the database is regularly updated with the latest measurements available. The data is public and

freely available through different interfaces ([WODselect web application](), THREDDS, HTTPS, FTP) in ASCII, Comma Separated Value (CSV) and netCDF formats. WOD is maintained by the National Centers for Environmental Information (NCEI) of the United States' National Oceanic and Atmospheric Administration (NOAA).

The International Quality-controlled Ocean Database (IQuOD) is an effort by the oceanographic community "with the goal of producing the highest quality and complete single ocean profile repository along with (intelligent) metadata and assigned uncertainties". It includes subsurface ocean profiles of several variables, paying particular attention to temperature measurements. The database is prepared by NCEI, and it is freely available in netCDF format through multiple channels (THREDDS, HTTPS, FTP).

Argovis is a REST API and web application hosted at the University of Colorado, Boulder (United States) and serving profile data in JSON format from the Argo program, and from ship and drifters missions.

The Ocean Data Platform by the non-profit HUB Ocean is among the youngest projects in the community, and it allows to find and access datasets from a catalog. The user can interact with the platform through different interfaces (SDK, REST API, OQS, JupyterHub workspaces), loading the datasets in tabular format.

The above efforts serve the data in ASCII, CSV, netCDF, or JSON formats. Generally speaking, netCDF is a binary format that offers the advantage to be compact and efficient when dealing with multidimensional data, while the others have the advantage of being human-readable (but can be very inefficient for large datasets). None of them is optimized for cloud object storage, although efforts are ongoing for netCDF (see Zarr and Icechunk). For this reason, Parquet has been drawing more attention from the earth sciences community recently: Parquet is a cloud-optimized binary format for tidy and large data (i.e. large tables) that is language-agnostic. It is widely used in the data science and corporate worlds, and the software ecosystem around it is sound and still growing. An overview of the characteristics of each format is in Table 1. We chose Parquet as the target format because it is optimized for cloud storage and cloud computing, its mature ecosystem includes packages in multiple coding languages to access it (Python, Julia, MATLAB, web technologies, etc.), and because the novel user is generally more familiar with tabular data and we want to make CrocoLake easily accessible from a technical standpoint too. Arguably, the main drawbacks of Parquet are that a large amount of workflows in ocean modeling are based on multidimensional data structures (not tabular) and that attaching attributes to the data can be clunky. At the same time, we see CrocoLake as one step in workflows that require fast access to point-based ocean observations, responding to necessities of data storage with fast access both on disk and the cloud, and not the end-all solution for ocean observations.

Table 1. Comparison between different common file formats for oceanographic datasets.

Feature Name	CSV	netCDF	Parquet	ASCII	JSON	Zarr + Icechunk
Cloud-Optimized	No	No	Yes	No	No	Yes
Structure	Tabular	Array-based	Tabular	Tabular	Hierar-chical	Array-based
Available tools in:						
Python	Yes	Yes	Yes	Yes	Yes	Yes
Julia	Yes	Yes	Yes	Yes	Yes	Zarr only
MATLAB	Yes	Yes	Yes	Yes	Yes	Zarr only

Attributes descriptors	Dedicated columns, header	Dictionary, accessed with data	Dictionary, separate access from data	Dedicated column	Dedicated field in object	Dictionary, accessed with data

Another key difference between CrocoLakeTools and the other projects is that CrocoLakeTools is open-source and anyone can use it to build their own flavor of CrocoLake and contribute to it by adding converters to support new datasets.

Code architecture

Converters

The core task of CrocoLakeTools is to take one or more files from a dataset and convert them to parquet, ensuring that CrocoLake's schema is followed. This is achieved through the methods contained in the `Converter` class and its subclasses. While the conversion of all datasets requires some general functionality (e.g. renaming the original variables to the final schema), each conversion requires specific tools for the specific dataset (e.g. the map used to rename the variables). CrocoLakeTools then contains the `Converter` class, which contains the methods shared across datasets, and from which converter subclasses inherit and implement the specific needs of each dataset.

Workflow

Local mirrors

The first step in the workflow is to retrieve the original files (Figure 1). The original sources follow the format, schema, nomenclature and conventions defined by the individual project (mission, scientist, etc.) the generated them and are unaware of CrocoLake's workflow. Modules to download the original data are optional. They should inherit from the `Downloader` class and be called downloader<DatasetName> (e.g. `downloaderArgoGDAC`). At the time of writing CrocoLakeTools is released with a downloader to build a local mirror of the Argo GDAC, and we hope to support more in the future. Whether a downloader module exists or the user downloads the data themselves, the original data is stored on disk and this is the starting point for the converter.

Parquet datasets

The second step is to convert the data to parquet, and finally merge the datasets into CrocoLake (Figure 2). The core of `CrocoLakeTools` are the modules in the `Converter` class and its subclasses. Each original dataset has its own subclass called converter<DatasetName>, e.g. `converterGLODAP`; further specifiers can be added as necessary (e.g. at this time there a few different converters for Argo data to prepare different datasets). The need for a dedicated converter for each project despite the usage of common data formats (e.g. netCDF, CSV) is due to differences in the schema, e.g. variable names, units, etc. Depending on the dataset, multiple converters can be applied. For example, to create CrocoLake, Argo data goes through two converters: 1. `converterArgoGDAC`, which converts the original Argo GDAC preserving most of its original conventions; 2. `converterArgoQC`, which takes the output of the previous step and applies some filtering based on Argo's QC flags and makes the data conforming to CrocoLake's schema.

CrocoLake

CrocoLake is one parquet dataset that contains each converted dataset merged together. This can be achieved with the script merge_crocolake.py. The script first creates a directory containing symbolic links to each converted dataset. It then uses the submodule CrocoLakeLoader to seamlessly load all the converted datasets into memory as one dask dataframe with a uniform schema and merge them into CrocoLake and stores it back to disk.

CrocoLake can be accessed with several programming languages with just a few lines of codes: the submodules CrocoLake-Python, CrocoLake-Matlab, and CrocoLake-Julia contain tools and examples in some languages.
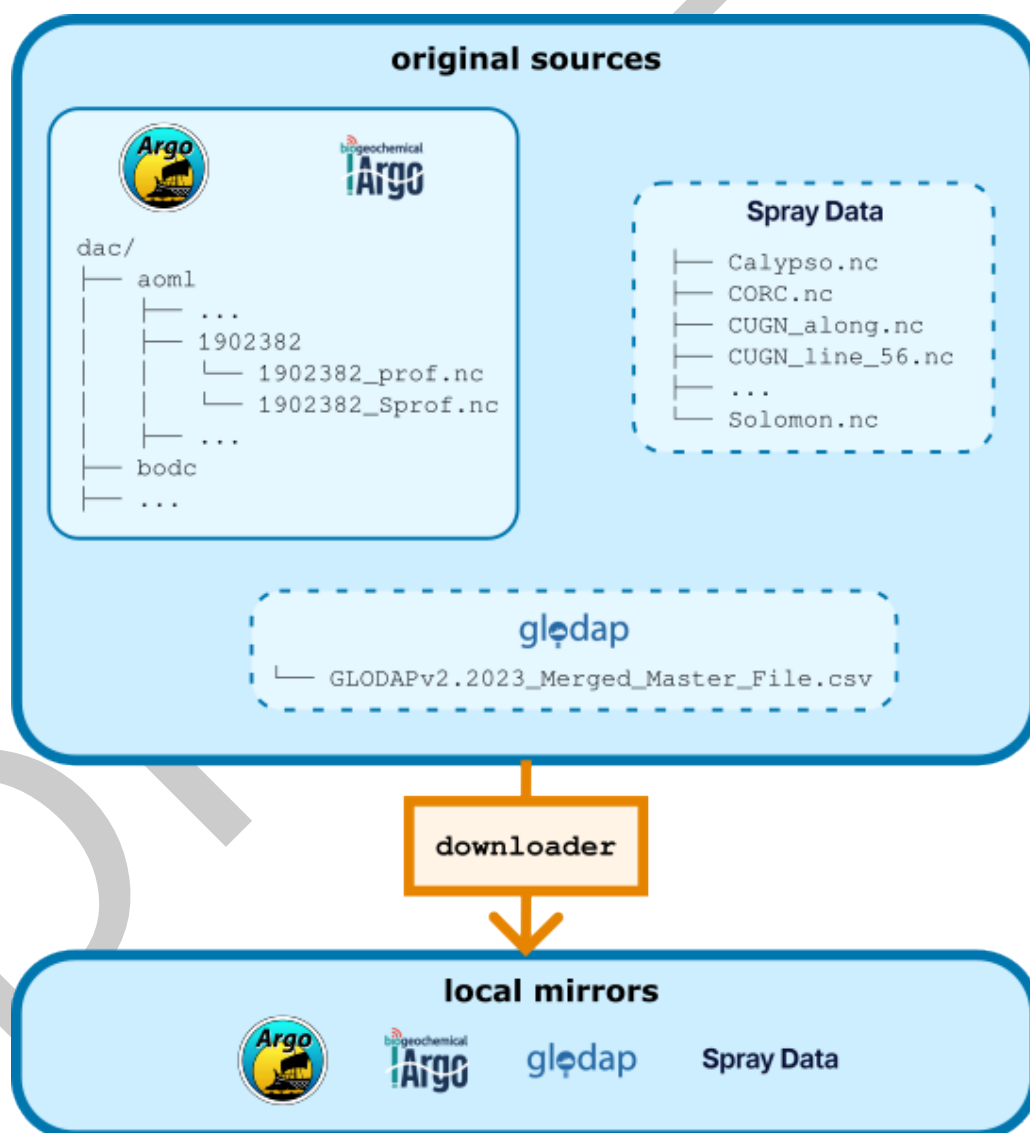

Figure 1: CrocoLake's workflow: 'downloader's. 'CrocoLakeTools' is set up to host modules that are dedicated to download the desired datasets from the web. It currently supports the download only of Argo data (solid line subset), and other datasets require the user to download them manually (dashed border subsets).

Figure 2: CrocoLake's workflowL 'converter's. 'converter's read the data in their original format, transform it following CrocoLake's conventions, converts it to parquet, and stores it back to disk. Each dataset is converted to its own parquet version. Thanks to the submodule 'CrocoLakeLoader', multiple parquet datasets are merged into a uniform dataframe which is save to disk as CrocoLake. For each dataset, a version containing only physical variables ('PHY') or also biogeochemical variables ('BGC') can be generated.

Schema

The nomenclature, units and data types are generally based on Argo's (https://vo-cab.nerc.ac.uk/collection/R03/current/). For CrocoLake's variables that are not present in the Argo program, we provide new names mantaining consistency with Argo's style.

Profile numbering

Ocean data is often accessed by profiles and we provide this functionality for CrocoLake too: the user can retrieve the profiles through the CYCLE_NUMBER variable, which is unique and progressive for each PLATFORM_NUMBER of each subdataset (DB_NAME). As each original product uses its own conventions, the CYCLE_NUMBER of some datasets is generated ad hoc during their conversion if no obvious match with CYCLE_NUMBER exists. The procedure for each dataset is detailed in the online documentation.

Quality control

CrocoLake contains only quality-controlled measurements. We rely exclusively on the quality-controls performed by the data providers and at the time we do not perform any control ourselves (although this might change in the future). Each parameter <PARAM> has a corresponding <PARAM>_QC flag that is generally set to 1 to indicate that the data is reliable. For Argo measurements, the original QC value is preserved, and only measurements with QC values of 1, 2, 5, and 8 considered.

Measurement errors

Each parameter <PARAM> has a corresponding <PARAM>_ERROR that indicates a measurement's error as provided in the original dataset. When no error is provided, <PARAM>_ERROR is set to null.

Documentation and updates

Documentation is available at (https://crocolakedocs.readthedocs.io/en/latest/index.html). It describes the specifics of each dataset (e.g. what quality-control filter we apply to each dataset, the procedure to generate the profile numbers, etc.), and it is updated every time a new feature is available.

Citation

If you use CrocoLakeTools and/or CrocoLake, do not limit yourself to citing this manuscript and remember to cite the datasets that you have used as indicated in the documentation. For example, if your work relies on Argo measurements, acknowledge Argo (Wong et al., 2020). This is important both for the maintainers of each product to track their impact and to acknowledge their efforts that made your work possible.

Acknowledgements

References

Wong, A. P., Wijffels, S. E., Riser, S. C., Pouliquen, S., Hosoda, S., Roemmich, D., Gilson, J., Johnson, G. C., Martini, K., Murphy, D. J., & others. (2020). Argo data 1999–2019: Two million temperature-salinity profiles and subsurface velocity observations from a global array of profiling floats. *Frontiers in Marine Science*, *7*, 700.